

## ML LAB 2

Explore and implement Linear regression algorithm in a given business scenario and comment on its efficiency and performance.

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import preprocessing
%matplotlib inline
```

```
In [5]: df=pd.read_csv("E:\DS\Datasets\winequalityN.csv")
```

```
In [7]: df.head(20)
```

Out[7]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	a
0	white	7.0	0.27	0.36	20.70	0.045	45.0	170.0	1.0010	3.00	0.45	
1	white	6.3	0.30	0.34	1.60	0.049	14.0	132.0	0.9940	3.30	0.49	
2	white	8.1	0.28	0.40	6.90	0.050	30.0	97.0	0.9951	3.26	0.44	
3	white	7.2	0.23	0.32	8.50	0.058	47.0	186.0	0.9956	3.19	0.40	
4	white	7.2	0.23	0.32	8.50	0.058	47.0	186.0	0.9956	3.19	0.40	
5	white	8.1	0.28	0.40	6.90	0.050	30.0	97.0	0.9951	3.26	0.44	
6	white	6.2	0.32	0.16	7.00	0.045	30.0	136.0	0.9949	3.18	0.47	
7	white	7.0	0.27	0.36	20.70	0.045	45.0	170.0	1.0010	3.00	0.45	
8	white	6.3	0.30	0.34	1.60	0.049	14.0	132.0	0.9940	3.30	0.49	
9	white	8.1	0.22	0.43	1.50	0.044	28.0	129.0	0.9938	3.22	0.45	
10	white	8.1	0.27	0.41	1.45	0.033	11.0	63.0	0.9908	2.99	0.56	
11	white	8.6	0.23	0.40	4.20	0.035	17.0	109.0	0.9947	3.14	0.53	
12	white	7.9	0.18	0.37	1.20	0.040	16.0	75.0	0.9920	3.18	0.63	
13	white	6.6	0.16	0.40	1.50	0.044	48.0	143.0	0.9912	3.54	0.52	
14	white	8.3	0.42	0.62	19.25	0.040	41.0	172.0	1.0002	2.98	0.67	
15	white	6.6	0.17	0.38	1.50	0.032	28.0	112.0	0.9914	3.25	0.55	
16	white	6.3	0.48	0.04	1.10	0.046	30.0	99.0	0.9928	3.24	0.36	
17	white	NaN	0.66	0.48	1.20	0.029	29.0	75.0	0.9892	3.33	0.39	
18	white	7.4	0.34	0.42	1.10	0.033	17.0	171.0	0.9917	3.12	0.53	
19	white	6.5	0.31	0.14	7.50	0.044	34.0	133.0	0.9955	3.22	0.50	

```
In [12]: df.columns
```

```
Out[12]: Index(['type', 'fixed acidity', 'volatile acidity', 'citric acid',  
              'residual sugar', 'chlorides', 'free sulfur dioxide',  
              'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol',  
              'quality'],  
             dtype='object')
```

```
In [13]: df.shape
```

```
Out[13]: (6497, 13)
```

```
In [14]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 6497 entries, 0 to 6496  
Data columns (total 13 columns):  
type                6497 non-null object  
fixed acidity       6487 non-null float64  
volatile acidity    6489 non-null float64  
citric acid         6494 non-null float64  
residual sugar      6495 non-null float64  
chlorides           6495 non-null float64  
free sulfur dioxide 6497 non-null float64  
total sulfur dioxide 6497 non-null float64  
density             6497 non-null float64  
pH                  6488 non-null float64  
sulphates           6493 non-null float64  
alcohol             6497 non-null float64  
quality             6497 non-null int64  
dtypes: float64(11), int64(1), object(1)  
memory usage: 659.9+ KB  
None
```

```
In [15]: df.isna().sum()
```

```
Out[15]: type                0  
fixed acidity              10  
volatile acidity           8  
citric acid                 3  
residual sugar             2  
chlorides                   2  
free sulfur dioxide         0  
total sulfur dioxide        0  
density                     0  
pH                           9  
sulphates                   4  
alcohol                     0  
quality                     0  
dtype: int64
```


```
In [16]: df=df.fillna(df.mean())
```

In [17]: df.describe()


Out[17]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000	6497.000000
mean	7.216579	0.339691	0.318722	5.444326	0.056042	30.525319	115.744574
std	1.295751	0.164548	0.145231	4.757392	0.035031	17.749400	56.521855
min	3.800000	0.080000	0.000000	0.600000	0.009000	1.000000	6.000000
25%	6.400000	0.230000	0.250000	1.800000	0.038000	17.000000	77.000000
50%	7.000000	0.290000	0.310000	3.000000	0.047000	29.000000	118.000000
75%	7.700000	0.400000	0.390000	8.100000	0.065000	41.000000	156.000000
max	15.900000	1.580000	1.660000	65.800000	0.611000	289.000000	440.000000


```
In [18]: import seaborn as sns
sns.set(rc={'figure.figsize':(10,8)})
corr = df.corr()
sns.heatmap(corr,
            xticklabels=corr.columns.values,
            yticklabels=corr.columns.values)
plt.show()
```



```
In [73]: plt.scatter("alcohol", "quality", data=df)  
plt.show()
```



```
In [19]: sns.boxplot(x="type",y="quality",data=df, palette="dark")  
plt.show()
```



```
In [20]: df=df[df.columns.drop('type')]
```

In [21]: `df.head(5)`

Out[21]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	0.45	8.8
1	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	0.49	9.5
2	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	0.44	10.1
3	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9
4	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9

In [22]: `print(df.nunique())`

```
fixed acidity      107
volatile acidity  188
citric acid        90
residual sugar    317
chlorides         215
free sulfur dioxide 135
total sulfur dioxide 276
density          998
pH               109
sulphates        112
alcohol          111
quality          7
dtype: int64
```

In [23]: `from sklearn.model_selection import train_test_split`  
`training, testing =train_test_split(df, test_size= 0.30, random_state=24)`

In [24]: `training.shape`

Out[24]: (4547, 12)

In [25]: `testing.shape`

Out[25]: (1950, 12)

In [28]: `X = training['alcohol']`

In [29]: `X.shape`

Out[29]: (4547,)

In [30]: `x= np.array(X)`

In [31]: `x = x.reshape(4547,1)`

```
In [32]: x.shape
```

```
Out[32]: (4547, 1)
```

```
In [33]: Y = training['quality']
```

```
In [34]: Y.shape
```

```
Out[34]: (4547,)
```

```
In [35]: Y= np.array(Y)
```

```
In [36]: y = Y.reshape(4547,1)
```

```
In [37]: y.shape
```

```
Out[37]: (4547, 1)
```

```
In [38]: from sklearn.linear_model import LinearRegression  
lr= LinearRegression()  
model=lr.fit(x, y)
```

```
In [39]: print(model)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [50]: print(model.coef_[0][0]) ## Printing the coefficients  
print(model.intercept_[0]) ### printing the Intercept term  
  
print("The linear model is: Y = {:.5} + {:.5}X".format(model.intercept_[0], model  
0.32546629798314014  
2.4029999180573034  
The linear model is: Y = 2.403 + 0.32547X
```

```
In [52]: X_test=testing['alcohol']
```

```
In [53]: X_test.shape
```

```
Out[53]: (1950,)
```

```
In [82]: X_test = X_test.reshape(1950,1)
```

```
In [83]: X_test.shape
```

```
Out[83]: (1950, 1)
```

```
In [84]: Y_test=testing['quality']
```

```
In [85]: Y_test.shape
```

```
Out[85]: (1950,)
```



In [86]:

```
Y_test = Y_test.reshape(1950,1)
```

C:\Users\PRANAV\Anaconda3\lib\site-packages\ipykernel\_launcher.py:2: FutureWarning: reshape is deprecated and will raise in a subsequent release. Please use .values.reshape(...) instead

In [88]: `Y_test.shape`

Out[88]: (1950, 1)

In [94]: `Y_test`Out[94]: 

```
array([[5],
       [4],
       [5],
       ...,
       [5],
       [7],
       [5]], dtype=int64)
```

In [89]: `Y_pred = lr.predict(X_test)`In [90]: `Y_pred`Out[90]: 

```
array([[5.72275616],
       [5.52747638],
       [5.3321966 ],
       ...,
       [5.56002301],
       [5.75530279],
       [5.46238312]])
```

In [91]: `from sklearn.metrics import mean_squared_error`In [92]: `LR_score= mean_squared_error(Y_test,Y_pred)`In [93]: `LR_score`

Out[93]: 0.610874859296884

## Interpretation:

The wine quality has been predicted using Linear Regression, with LR score of 61%

In [ ]:

